

```
In [1]: # Importing Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from sklearn.svm import SVR
from sklearn.metrics import r2_score, mean_squared_error, mean_absolute_error
```

1. Loading and Preprocessing

```
In [3]: # Load the dataset

df = pd.read_csv('CarPrice_Assignment.csv')
```

```
In [5]: # Preview the dataset

print("Dataset preview:")
df.head()
```

Dataset preview:

```
Out[5]:
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drive
0	1	3	alfa-romero giulia	gas	std	two	convertible	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	
3	4	2	audi 100 ls	gas	std	four	sedan	
4	5	2	audi 100ls	gas	std	four	sedan	

5 rows × 26 columns



```
In [7]: # Check for missing values

print("\nMissing values in each column:")
df.isnull().sum()
```

Missing values in each column:

```
Out[7]: car_ID      0
        symboling   0
        CarName     0
        fueltype    0
        aspiration   0
        doornumber   0
        carbody      0
        drivewheel   0
        enginelocation 0
        wheelbase    0
        carlength    0
        carwidth     0
        carheight    0
        curbweight   0
        enginetype   0
        cylindernumber 0
        enginesize    0
        fuelsystem   0
        boreratio    0
        stroke       0
        compressionratio 0
        horsepower   0
        peakrpm      0
        citympg      0
        highwaympg   0
        price        0
        dtype: int64
```

```
In [9]: # Drop irrelevant columns (if any)

        # Drop car_id if it doesn't contribute to regression
        if 'car_ID' in df.columns:
            df.drop(columns=['car_ID'], inplace=True)
```

```
In [11]: # Encoding categorical variables (if any)

        df = pd.get_dummies(df, drop_first=True)
```

```
In [13]: # Feature scaling
        scaler = StandardScaler()
        X = df.drop('price', axis=1)
        y = df['price']
        X_scaled = scaler.fit_transform(X)
```

```
In [15]: # Splitting dataset into training and testing sets
        X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, test_size=0.2,
```

2. Model Implementation

```
In [17]: # Initialize regression models

        models = {
            "Linear Regression": LinearRegression(),
            "Decision Tree": DecisionTreeRegressor(random_state=42),
            "Random Forest": RandomForestRegressor(random_state=42),
            "Gradient Boosting": GradientBoostingRegressor(random_state=42),
```

```
"Support Vector Regressor": SVR()
}
```

```
In [19]: # Train and evaluate each model
model_performance = []
for name, model in models.items():
    print(f"\nTraining {name}...")
    model.fit(X_train, y_train)
    y_pred = model.predict(X_test)

    # Evaluate performance
    r2 = r2_score(y_test, y_pred)
    mse = mean_squared_error(y_test, y_pred)
    mae = mean_absolute_error(y_test, y_pred)
    model_performance.append((name, r2, mse, mae))

    print(f"{name} Performance:")
    print(f"R-squared: {r2:.4f}")
    print(f"Mean Squared Error: {mse:.4f}")
    print(f"Mean Absolute Error: {mae:.4f}")
```

Training Linear Regression...

Linear Regression Performance:

R-squared: -31650981667885770354982912.0000

Mean Squared Error: 2498655757651236794056222502289408.0000

Mean Absolute Error: 30899709266986604.0000

Training Decision Tree...

Decision Tree Performance:

R-squared: 0.8559

Mean Squared Error: 11376015.6135

Mean Absolute Error: 2200.1423

Training Random Forest...

Random Forest Performance:

R-squared: 0.9533

Mean Squared Error: 3682803.2185

Mean Absolute Error: 1367.3156

Training Gradient Boosting...

Gradient Boosting Performance:

R-squared: 0.9308

Mean Squared Error: 5463056.2200

Mean Absolute Error: 1696.8629

Training Support Vector Regressor...

Support Vector Regressor Performance:

R-squared: -0.1021

Mean Squared Error: 87001508.9811

Mean Absolute Error: 5707.0130

3. Model Evaluation

```
In [21]: # Create a DataFrame to compare results

performance_df = pd.DataFrame(model_performance, columns=["Model", "R-squared",
performance_df = performance_df.sort_values(by="R-squared", ascending=False)
```

```
print("\nModel Performance Comparison:")
print(performance_df)
```

Model Performance Comparison:

	Model	R-squared	MSE	MAE
2	Random Forest	9.533492e-01	3.682803e+06	1.367316e+03
3	Gradient Boosting	9.307984e-01	5.463056e+06	1.696863e+03
1	Decision Tree	8.558977e-01	1.137602e+07	2.200142e+03
4	Support Vector Regressor	-1.020658e-01	8.700151e+07	5.707013e+03
0	Linear Regression	-3.165098e+25	2.498656e+33	3.089971e+16

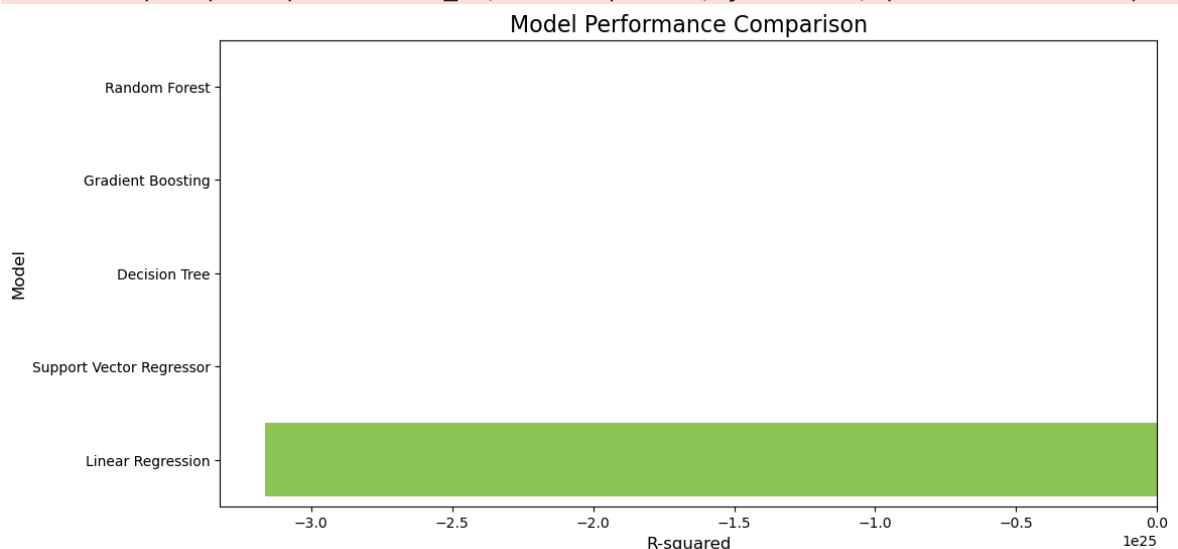
In [34]: *# Visualization*

```
plt.figure(figsize=(12, 6))
sns.barplot(data=performance_df, x="R-squared", y="Model", palette="viridis")
plt.title("Model Performance Comparison", fontsize=16)
plt.xlabel("R-squared", fontsize=12)
plt.ylabel("Model", fontsize=12)
plt.show()
```

C:\Users\vayal\AppData\Local\Temp\ipykernel_10440\2079656924.py:4: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(data=performance_df, x="R-squared", y="Model", palette="viridis")
```



In [23]: *# Identify best and worst models*

```
best_model_name = performance_df.iloc[0]['Model']
print(f"\nBest Performing Model: {best_model_name}")
```

Best Performing Model: Random Forest

4. Feature Importance Analysis

In [25]: *# For tree-based models, extract feature importances*

```
if best_model_name in ["Random Forest", "Gradient Boosting", "Decision Tree"]:
    best_model = models[best_model_name]
    feature_importances = pd.DataFrame({
        'Feature': X.columns,
```

```

'Importance': best_model.feature_importances_
}).sort_values(by='Importance', ascending=False)

print("\nFeature Importance:")
print(feature_importances)

# Plot feature importances
plt.figure(figsize=(12, 6))
sns.barplot(data=feature_importances.head(10), x="Importance", y="Feature",
plt.title("Top 10 Feature Importances", fontsize=16)
plt.xlabel("Importance", fontsize=12)
plt.ylabel("Feature", fontsize=12)
plt.show()

```

Feature Importance:

	Feature	Importance
6	enginesize	0.540808
5	curbweight	0.294927
13	highwaympg	0.043387
10	horsepower	0.038275
3	carwidth	0.013358
..
142	CarName_volkswagen rabbit	0.000000
91	CarName_nissan note	0.000000
118	CarName_subaru baja	0.000000
55	CarName_honda civic 1500 gl	0.000000
147	CarName_volkswagen rabbit	0.000000

[189 rows x 2 columns]


C:\Users\vayal\AppData\Local\Temp\ipykernel_5016\2145102746.py:14: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v 0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```

sns.barplot(data=feature_importances.head(10), x="Importance", y="Feature", palette="viridis")

```



5. Hyperparameter Tuning

```
In [28]: # Hyperparameter tuning for Random Forest
if best_model_name == "Random Forest":
    param_grid = {
        'n_estimators': [100, 200, 300],
        'max_depth': [None, 10, 20, 30],
        'min_samples_split': [2, 5, 10],
        'min_samples_leaf': [1, 2, 4]
    }
    grid_search = GridSearchCV(RandomForestRegressor(random_state=42), param_grid)
    grid_search.fit(X_train, y_train)

    print("\nBest Hyperparameters for Random Forest:")
    print(grid_search.best_params_)
```

Best Hyperparameters for Random Forest:

{'max_depth': None, 'min_samples_leaf': 2, 'min_samples_split': 2, 'n_estimators': 300}

```
In [32]: # Evaluate the tuned model
tuned_model = grid_search.best_estimator_
y_tuned_pred = tuned_model.predict(X_test)
tuned_r2 = r2_score(y_test, y_tuned_pred)
print(f"\nTuned Model R-squared: {tuned_r2:.4f}")
```

Tuned Model R-squared: 0.9544

```
In [36]: # Save results for submission
performance_df.to_csv('Car_Price_Model_Performance.csv', index=False)
```

In []: